





Pod Execution Model

How pods ship working software as fast as possible



3

Why Pods Exist


Pods compress and collapse the delivery timeline.

- Small cross-functional strike teams
- Iterate first, not spec first
- Collapse roles, remove translation layers
- AI fills the gaps that used to require more people

The model is already proven with kCapture. Calendar shipping fast makes it spread.

Continuous Iteration

Every delivery strengthens the next. The system learns as it runs.



Execution becomes predictable when learning never stops.

01 Discover

Synthesize stakeholder signals, system context, and competitive inputs to define a differentiated solution direction.

DELIVERABLE: SIGNED-OFF DOCUMENTATION + PROMPT PACKAGE

02 Iterate

Iterate through working Canvas-driven iterations with same-day feedback and pre-code stakeholder alignment. Then build, integrate, and continuously test against acceptance criteria.

DELIVERABLE: VALIDATED ITERATION READY FOR PRODUCTION SCAFFOLDING

03 Launch

Ship to production users, run retrospective, capture learnings, and feed evidence back into the next Discover cycle.

DELIVERABLE: RELEASE + RETRO INTELLIGENCE REPORT

5

Pod Structure

Execution cores, not committees. Two operators drive the pod. Everyone else advises on demand.

CORE OPERATORS

 Engineer

 Champion

AVAILABLE SUPPORT · ON DEMAND

UX

QA

PM

CS

- Engineer and Champion drive the pod. Everyone else advises.
- AI is the force multiplier. Every person operates at 3-5x output.

6

Inputs Replace Deliverables

The old model passed documents between roles. The new model uses AI to generate working inputs into a shared iteration.

The Old Way

- ✗ PM writes PRD over days, hands to Design
- ✗ Designer spends weeks in Figma, hands to Eng
- ✗ Engineer builds from spec, hands to QA
- ✗ Each step adds weeks and loses context

The New Way: Roles Collapse

- ✓ Anyone with VS Code can draft, design, or build
- ✓ Champion drafts the problem frame AND helps build a clickable working version
- ✓ Designer generates supporting documentation AND ships clickable UI/visuals.
- ✓ Engineer scaffolds the app while QA agents, kAutomate, and automation generate structured validation coverage
- ✓ No handoffs. Everyone contributes everywhere.

AI removes the specialization and man-power needed, teams of 10 become teams of 2 or 3.

7

Industry Inputs

Feed AI all available context: customer signals, competitive products, internal presentations, articles, and domain expertise. No multi-week spec cycles.

Gather Context

- Feed meeting transcriptions directly into ChatGPT
- Paste Teams conversations, support tickets, customer interviews
- Output: structured summary of what's broken and why it matters

ChatGPT

Teams

Copilot

Draft Documentation

- Ask AI to generate minimal documentation from all gathered context
- Scope to core user flows. Cut everything else.
- Output: 1-page with problem, scope, constraints, and "done" definition

ChatGPT

Copilot

Define Done

- AI drafts testable acceptance criteria from the documentation
- Lock what "shipped" means before anyone builds. Cut ambiguity early
- Output: clear definition of done the pod can ship against

ChatGPT

Copilot

What used to require heavy manual synthesis now takes a morning with AI. Gather context. Draft documentation.

8

Designer Inputs

UI/UX accelerates refinement after a working iteration, not by defining pixel-perfect specs upfront. Build clickable working versions in a couple hours, not weeks.

Shape the Brief

- Take draft documentation and pressure-test it with UX thinking
- Feed user research and competitor screenshots into ChatGPT to surface gaps
- Output: refined brief with interaction requirements

ChatGPT

Copilot

Build in Code

- Open VS Code + an approved model. Describe the UI in plain language
- Full page with states, navigation, real copy. Working iteration, not static mockups
- Output: clickable iteration the pod can test same day

VS Code

Opus 4.6

Codex

Iterate Live

- Tell AI "make the sidebar narrower" or "add dark mode." See it instantly
- Refine the iteration alongside the engineer in the same codebase
- Designer becomes a creative director of AI output, not a pixel pusher

VS Code

Opus 4.6

A designer with VS Code can produce a clickable artifact in a single iteration session.

9

Engineer Inputs

VS Code + an approved model scaffolds the app. Engineer directs, integrates, and ensures quality.

Scaffold the Real App

- Take iterations from PM and designer. Use AI as a pair partner, pointed at our codebase.
- Scaffold production code from iteration patterns. Bridge the gap from validation to production
- Primary user flow running end-to-end in the real app within 48 hours

VS Code

Opus 4.6

Codex

Ship and Iterate

- "Add form validation" / "add offline support."
Done in minutes, not sprints
- AI generates tests, API contracts, data models, and component variations
- Engineer's role shifts from writing code to directing AI and ensuring quality

VS Code

Opus 4.6

Codex

Own the Architecture

- Map integration surfaces: APIs, auth, data sources
- Use AI as a peer reviewer: surface risks, challenge assumptions, catch blind spots
- Deployable structure from day 1: CI, env config, feature flags

ChatGPT

VS Code

Engineers now orchestrate AI while continuous feedback replaces translation cycles and drives rapid iteration.

10

Stakeholder Inputs

QA supports scope definition when needed. Stakeholders remove friction, not approve output.

QA Advisory Support

- QA supports scope definition when needed, not embedded as standing labor
- Validates the work and shapes requirements through what breaks
- Regression and release validation still exists downstream
- Uses AI to generate test scenarios, edge cases, and regression checklists from the evolving iteration

ChatGPT

VS Code

Automation-Driven Coverage

- Automation generates QA coverage, people support edge cases
- Engineer uses AI to generate automated tests alongside every feature
- Designer tests interactions in the working iteration, not in Figma comments

VS Code

Opus 4.6

Stakeholder Alignment

- Strategic framing: why this matters right now
- Fast yes/no on scope questions. No committees
- Remove organizational friction. Clear the path so the pod can ship

QA is consultative, not gating. Automation drives coverage. Stakeholders clear the path.

11

Artifact Expectations

Every artifact should accelerate testing or iteration. If it doesn't, cut it.

What Pods Produce

- ✓ **Rapid Project Plan** is the live execution artifact pods iterate against
- ✓ Clickable iterations (not static mockups)
- ✓ AI-generated test checklists and edge cases
- ✓ Working demo for stakeholder feedback
- ✓ Short decision notes ("we chose X because Y")
- ✓ Advisors review **Rapid Project Plan** asynchronously when needed

What Pods Don't Produce

- ✗ Multi-page spec documents
- ✗ Polished Figma handoff files
- ✗ Status decks or approval presentations
- ✗ Documentation before the thing works
- ✗ Anything that requires a meeting to review

The artifact is the iteration. Everything else is a note that helps the iteration get better.

12

Rapid Project Plan: First Artifact

The required entry artifact for every pod before work begins.

Problem

What problem are we solving right now?

Primary User

Who runs the workflow?

Definition of Done

Operational · Usable · Demonstrable

Owner

Champion responsible for execution

Scope Boundary

What is explicitly out of scope?

Test Workflow

Exact capture-to-review flow

Success Signal

What proves value quickly?

This is the required entry artifact for every pod before work begins.

Rapid Project Plan

This template is actively being modified and standardized. Early example below.

Miami Dade Fast Track

TL;DR

Stabilize File Manager performance for large folder structures by migrating folder permissions to the new domain permission model, then deliver bulk permission management to eliminate folder-by-folder edits.

Success = large projects load without degradation, admins modify permissions across thousands of folders efficiently.

What We're Shipping

- **Backend Architecture:** Migrate to new domain permission model
- **Performance Stabilization:** Eliminate deep folder permission crawl
- **UI Enhancement:** Bulk permission management workflow

Problem Breakdown

Performance

- Deep folder trees degrade performance
- Thousands of folders per project
- API workarounds do not scale

Usability

- Manual edits per folder required
- Miami Dade: ~4000 folders, ~75 groups
- Current UI not viable at enterprise scale

Goal

Ensure File Manager supports enterprise-scale folder structures without performance degradation while enabling efficient bulk permission management.

Scope — Phase 1: Performance

- Replace folder permission logic with new domain model
- Align folder architecture with partition-level permissions
- Validate performance on large data sets

Out of scope: usability enhancements before performance stabilization

Scope — Phase 2: Usability

- Introduce bulk permission editing workflow
- Enable cross-folder updates in a single action
- Ensure solution benefits all customers, not just Miami Dade

Core Technical Decisions

- Adopt onboarding rewrite permission model for folders
- Avoid API-based mass permission mutation
- Performance validation before UI improvements begin

Constraints & Trade-offs

- Euro engineering team owns implementation
- Euro team is resource constrained
- Performance must be solved before usability
- High-visibility customers involved
- No external engineering support for Euro team

Phased Plan

Phase 1: Define migration approach · Assign engineering ownership · Validate performance on large projects

Phase 2: Design bulk permission UX · Implement efficient update workflow · Test across Miami Dade and GSA

Phase 3: Generalize rollout across enterprise · Monitor performance and feedback

Stakeholders

- Engineering Architecture: Colin
- Implementation: Vlad and Euro Team
- API Context: Mike Richie
- Customer Success: Kimberly Day
- Partner Enablement: John Quigley

14

Pod Execution Vocabulary

Prevents terminology drift across pods.

Pod

A small execution unit built around two core operators and on-demand advisors. Not a traditional team. A focused delivery core designed to move at the speed of decisions.

Advisor

UX, PM, QA, Platform, or any role with relevant expertise. Contributes when the pod needs it. Supports progress rather than gating it. Not standing headcount on any given pod.

Iteration

A working, clickable representation of the workflow. Each one is better than the last. The point is to get something in front of people fast, gather feedback, and improve from there.

Iteration Loop

Discover, Iterate, Validate, Ship, Repeat. The continuous execution cycle pods operate within. The tempo adapts to the problem. Each pass tends to get faster as patterns emerge and become reusable.

Demo

A working demonstration inside the real system, not slides or screenshots. Stakeholder feedback happens here. If it can't be shown working, it likely isn't ready yet.

Champion

The person closest to the problem who drives scope, removes blockers, and makes decisions. Accountable for what ships and whether it solves the right thing. Could be a PM, could be anyone.

Rapid Project Plan

A short execution source of truth. Defines the problem, scope boundary, definition of done, and owner. Intended to be a living document that evolves as the pod learns, not a locked spec.

Inputs

Context gathered from any available source and fed into AI to produce working artifacts. Meeting transcripts, support tickets, competitive screenshots, customer interviews. Replaces sequential deliverables between roles.

Done

Operational and usable. A real user can complete the core workflow. Does not mean polished or fully documented. The goal is shipped capability that can be demonstrated and improved.

Shared language eliminates misalignment. Every pod uses these terms the same way.

15

Iteration Speed

Continuous operational progress beats delayed polish.

- **Visible progress:** The working iteration advances continuously
- **Stakeholder reaction:** Feedback happens inside the real system, not in review meetings
- **Cross-pod learning:** Successful patterns spread immediately across teams

Execution is the strategy. The pod that ships fastest teaches the org the most.

16

Definition of Done

Done means operational, testable, usable. Not polished, finalized, documented.

Done Means

- ✓ Operational: works end to end
- ✓ Testable: someone can try it
- ✓ Adoptable: user completes the core flow
- ✓ Demonstrable: showable in a working demo

Done Does Not Mean

- ✗ Pixel perfect
- ✗ Fully documented
- ✗ All edge cases handled

Pods ship capability, not perfection.

**"Teams no longer operate around a PRD.
They operate around an Iteration."**

18

kCapture Pod: Case Study

Floor-plan-anchored inspection capture. Pins hold location history over time. Mobile captures, desktop reviews.

POD COMPOSITION

ProCon

Champion

Colin

Engineer

Bridgette

UX Advisor

ALREADY SHIPPED

- Floor plans with persistent inspection pins
- Capture timelines per location (photo, video, 360, files)
- Live mobile 360 panorama with guided rotation
- Offline-tolerant upload queue with transfer center
- Side-by-side synchronized comparison across time

WHAT THE POD IS PROVING

- One user can capture, upload, and review without training
- Capture always binds to the correct pin
- Location history is clear on repeat visits
- Sequential pin workflow supports checklist-style inspections

Spatial inspection, not media storage. Capture attaches to place, not gallery.

[illegible]

20

Canvas Pod: Case Study

The first pod that proved the model. Canvas shipped as a real product using the execution framework before it was formalized.

POD COMPOSITION

Pritesh

Champion

Colin

Engineer

Sarah

UX Advisor

WHAT IT PROVED

- Small team, fast iteration, working product in weeks
- AI-assisted development accelerated every phase
- Stakeholder feedback happened inside the real iteration
- No spec-first pipeline. Problem → iterate → ship

Canvas is the pattern. These three pods scale it.

21

Three Pilot Pods

One execution model. Three proving environments. Each validates a different adoption surface.

Calendar Pod

Engineer: Colin

Champion: Steve Stankiewicz

Advisors: Denny, Bridgette

- Ship Calendar as a proving surface for rapid execution on real workflows
- Example pattern: Calendar for Submittals, Calendar for Inspections, Calendar for Closeout
- This is the proving ground. If it works, the model expands org-wide after May release

Strategic Apps Pod

Engineer: Vlado

Champions: Sakiba, Jonathan

Advisors: Brian Haines

- Strategic execution on AssetCentric Project Management and expansion surfaces around it
- Two app docs already exist as starting material for immediate execution
- Proves the model works for strategic platform evolution, not just greenfield builds

Forward Deployed Pod

Engineer: --

Champion: Jordan Hoff

Advisors: --

- Embed with owner customers. Visit, interview, understand the work first
- Discovery inside real workflows followed by immediate implementation
- Engineering support required before activation

Small teams, fast inputs. Deliver in execution cycles, not planning cycles.

22

Timeline

Definition to deployment in rapid cycles. Calendar is the pilot.

- **Now:** Staff the Forward Deployed pod (engineer + designer for Jordan). Confirm all three pod rosters
- **Next:** Education and alignment meeting with everyone in all three pods. Calendar pod kicks off immediately after
- **As Fast As Possible:** Calendar pod ships working product. Model validated. Team can articulate what inputs are actually needed
- **After May Release:** Broader rollout begins pending validation from initial three pods. Calendar pod becomes the template every new team follows

23

What Changes Now

Concrete behavior shifts starting immediately.

- No more weeks of polish before learning. Ship unfinished but operational
- PMs ship a working iteration the same day they write the documentation. Problem frames become working artifacts
- Designers move into VS Code. Layout ships as working HTML, not Figma files
- Engineers direct AI, not write every line. Scaffolding, iteration, testing all accelerated
- QA is everyone's job, embedded from day 1. No downstream handoff
- Pod execution coordination and artifact tracking owned centrally


APPENDIX

Pod Execution System

25

The Flow

From problem to product. Four phases. No fixed timeline. Scale to fit.



Each phase has a clear output. The cycle runs at any scale. Same discipline.

26

Who Drives Each Phase

Engineer and Champion are constant. All other roles participate as needed.

	DISCOVER	ITERATE	BUILD	LAUNCH & DEPLOY
Engineer	Tech assessment, integration mapping, env setup	Scaffold real app from iteration patterns	Core flow end-to-end, API integration, testing	Bug fixes, edge cases, regression, deploy
Champion	Scope lock, define done, set cadence	Stakeholder demo, capture feedback, hold scope	Mid-cycle check, kill scope creep, unblock	Release decision, retrospective, capture lessons
PM *	Meetings, synthesize documentation, lock scope	Build clickable version, validate with users	Refine scope, cut non-core, support testing	Soft launch coordination, feedback capture
Designer *	Research, competitor audit, pattern library	Refine iteration, add all states and flows	Polish interactions directly in codebase	Final interaction polish, launch visuals
QA *	Define test strategy from documentation	Test iteration, shape requirements through bugs	Continuous testing on real builds, edge cases	Regression pass, release sign-off


* Advisory roles — can be filled by engineer, champion, or on-demand advisor depending on pod composition

Roles overlap. Phases compress. Rows marked * are advisory functions, not required headcount.

27

Iteration Mechanisms

The system doesn't end at launch. Every cycle feeds the next.



After every release

Retro within 48 hours. What shipped, what slipped, what we learned.

Continuous feedback

Usage data flows directly into the next cycle. No separate research sprint.

Compounding speed

Patterns become reusable. The pod gets faster. Each cycle builds on the last.

Speed is earned, not assumed. The system compounds.

28

Local Validation Scales to Market Patterns

Clarifies how pods avoid one-off solutions.

- **Validate locally:** Pods validate workflows with real users immediately
- **Aggregate centrally:** Signals aggregate across pods centrally
- **Graduate patterns:** Repeated patterns graduate into platform capabilities
- **Insight through reuse:** Single-customer insight becomes market insight through reuse

One customer validates the workflow. Multiple customers confirm the pattern. The platform ships the capability.

29

PM: Old Way vs New Way

From spec author to decision scientist. The PM's value is judgment, not documentation.

The Old Way

- ✗ Spends weeks writing detailed PRDs before anyone builds
- ✗ Gathers requirements through meetings and documents them in slides
- ✗ Hands finished spec to design, waits for mockups, hands to engineering
- ✗ Success measured by spec completeness, not product learning
- ✗ Manages timelines and Jira tickets, not outcomes

The New Way

- ✓ Defines problem frame inside a live iteration on day one
- ✓ Feeds stakeholder signals, competitive context, and customer data directly into AI to synthesize direction
- ✓ Builds a working clickable version same day using Canvas and approved models
- ✓ Validates assumptions through real user interaction, not spec review
- ✓ Operates as a decision scientist: chooses which problems to solve, not how to document them

The PM who ships an iteration discovers what works. The PM who ships a PRD describes what might.

30

Designer: Old Way vs New Way

From pixel pusher to creative director of AI output. Design accelerates iteration, not precedes it.

The Old Way

- ✗ Spends weeks in Figma creating pixel-perfect mockups nobody has tested
- ✗ Delivers static screens with redline annotations for engineering to interpret
- ✗ Design-to-code handoff creates translation loss and rework cycles
- ✗ Feedback happens in comments on static files, disconnected from real behavior
- ✗ Waits for PRD before starting any visual work

The New Way

- ✓ Opens VS Code + an approved model and describes UI in plain language
- ✓ Produces working interactive iterations in hours, testing real states and navigation
- ✓ Directs and curates AI-generated output like a creative director, not a production artist
- ✓ Iterates live alongside engineering in the same codebase, no handoff
- ✓ Refines after a working iteration exists, not before one is built

A designer with an AI pair can now do work that once required a small team. Roles collapse into one workflow.

31

Engineer: Old Way vs New Way

From line-by-line coder to AI orchestrator. The engineer's value is architecture and judgment, not keystrokes.

The Old Way

- ✗ Waits for spec and mockups before writing a single line of code
- ✗ Builds from scratch, manually writing boilerplate, tests, and API contracts
- ✗ Sprints measured by story points and velocity, not working product
- ✗ Code review is the only quality gate, happens after completion
- ✗ Integration and deployment are separate phases with separate teams

The New Way

- ✓ Takes iteration from PM/Designer and scaffolds production code from patterns on day one
- ✓ Uses AI as a pair partner: generates tests, API contracts, data models, and component variations
- ✓ Primary user flow running end-to-end in the real app within 48 hours
- ✓ Continuous feedback replaces translation cycles. Iteration happens in minutes, not sprints
- ✓ Owns architecture, integration surfaces, and deployment from day one with CI, env config, and feature flags

The engineer who directs AI and ships working results outpaces the team that hand codes every step.

32

QA: Old Way vs New Way

From post-delivery gatekeeper to embedded quality partner. Validation is continuous, not sequential.

The Old Way


- ✗ Tests after engineering declares "done." Finds bugs after context is lost
- ✗ Writes manual test plans from static specifications
- ✗ QA is a phase: gated, scheduled, and siloed from development
- ✗ Regression testing is manual, slow, and incomplete
- ✗ "Hand off to QA" creates adversarial dynamic and delays release

The New Way


- ✓ Joins at pod kickoff. Tests the first iteration on day one
- ✓ QA agents, kAutomate, and AI generate test scenarios, edge cases, and regression checklists continuously
- ✓ Everyone is QA: PM validates assumptions, Designer tests interactions, Engineer generates automated tests
- ✓ Validates on working iterations, not specifications. Bugs caught in context, not after handoff
- ✓ QA supports pods and accelerates release, never gates it


There is no "hand off to QA." QA is already there. Validation is continuous, not a checkpoint.

HORSE AND CARRIAGE
The Old Way (The Carriage)



ROCKET SHIP
The New Way (The Rocket)





CURRENT DRAFT

March 26, 2026